


An example session (block width = length of the user turn)



2 Mine weak spots with a soft prompt

soft prompt turns 1–3 G (frozen) next word



gradients update P only

G's word-embedding space

- ★ F's word “temple”
- its neighbours
- ◆ G's mean guess (near: weigh more)
- ▶ push G away

keep entropy



spike spread

best of M

hardness per turn: t1, t3 = weak spots (G easy to push)

3 Patch G locally

base weights + LoRA (low-rank) on attention + MLP



weight = hardness

t1 t2 t3

F's replies

prompt discarded

trained frozen

4 Gate, serve, roll back

gate: both checks must pass

replies agree

query on topic

summary + last K turns full history

drift monitor

2 in a row

back to F