



1 Warm-Start

Early dialogue turns
(1 ... W)



Teacher replies
(cache and fix)


Warm-Start Pairs




Cache Teacher Replies

2 Soft-Prompt Mining


Input assembly (teacher-forced)



Mining Objective (teacher-forced student predictions vs. cached targets)




All Warm-Start Turns · Weights $\omega_t \propto r_t$



Optimize P · Freeze G · Anchoring Scores r_t

3 Local LoRA Distillation


All W warm-start turns
(anchoring weights $\omega_t \propto r_t$)



4 Gated Serving

Acceptance Checks

- Response Fidelity (acceptance batch vs. F)
- Query Locality (vs. session centroid)



Prompt-Free Serving · No Online LoRA Updates

Session timeline

